

Fast forecast reconciliation using linear models

Mahsa Ashouri

Institute of Service Science, National Tsing Hua University, Taiwan

Email: mahsa.ashouri@iss.nthu.edu.tw

Corresponding author

Rob J Hyndman

Monash University, Clayton VIC 3800, Australia

Email: rob.hyndman@monash.edu

Galit Shmueli

Institute of Service Science, National Tsing Hua University, Taiwan

Email: galit.shmueli@iss.nthu.edu.tw

22 November 2019

JEL classification: C10,C14,C22

Fast forecast reconciliation using linear models

Abstract

Forecasting hierarchical or grouped time series usually involves two steps: computing base forecasts and reconciling the forecasts. Base forecasts can be computed by popular time series forecasting methods such as Exponential Smoothing (ETS) and Autoregressive Integrated Moving Average (ARIMA) models. The reconciliation step is a linear process that adjusts the base forecasts to ensure they are coherent. However using ETS or ARIMA for base forecasts can be computationally challenging when there are a large number of series to forecast, as each model must be numerically optimized for each series. We propose a linear model that avoids this computational problem and handles the forecasting and reconciliation in a single step. The proposed method is very flexible in incorporating external data, handling missing values and model selection. We illustrate our approach using two datasets: monthly Australian domestic tourism and daily Wikipedia pageviews. We compare our approach to reconciliation using ETS and ARIMA, and show that our approach is much faster while providing similar levels of forecast accuracy.

Keywords: hierarchical forecasting, grouped forecasting, reconciling forecast, linear regression

1 Introduction

Modern data collection tools have dramatically increased the amount of available time series data. For example, the Internet of Things and point-of-sale scanning produce huge volumes of time series in a short period of time. Naturally, there is an interest in forecasting these time series, yet forecasting large collections of time series is computationally challenging.

1.1 Hierarchical and grouped time series

In many cases, these time series can be structured and disaggregated based on hierarchies or groups such as geographic location, product type, gender, etc. An example of hierarchical time series is sales in restaurant chains, which can be disaggregated into different stores and then different types of food or drinks. Figure 1 shows a schematic of such a hierarchical time series structure with three levels. The top level (level 0) is the total series, formed by aggregating all the bottom level series. In the middle level (level 1), series are aggregations of their own child

series; for instance, series A is the aggregation of AW and AX. Finally, the bottom level series (level 2), includes the most disaggregated series.




Figure 1: An example of a two level hierarchical structure.

Grouped time series involve more complicated aggregation structures compared to strictly hierarchical time series. To take the simplest example, suppose we have two grouping factors which are not nested: sex (Male/Female) and city (New York/San Francisco). The disaggregated series for each combination of sex and city can be combined to form city sub-totals, or sex sub-totals. These sub-totals can be combined to give the overall total. Both sub-totals are of interest.

We can think of such structures as hierarchical time series without a unique hierarchy. A schematic of this grouped time series structure is shown in Figure 2 with two grouping factors, each of two levels (A/B and C/D). The series in this structure can be split first into groups A and B and then subdivided further into C and D (left side), or split first into C and D and then subdivided into A and B (right side). The final disaggregation is identical in both cases, but the middle level aggregates are different.

We use the same notation (following Hyndman & Athanasopoulos 2018) for both hierarchical and grouped time series. We denote the total series at time t by y_t , and the series at node Z (subaggregation level Z) and time t by $y_{Z,t}$. For describing the relationships between series, we use an $N \times M$ matrix, called the ‘summing matrix’, denoted by S , in which N is the overall number of nodes and M is the number of bottom level nodes. For example in Figure 1, $N = 7$ and $M = 4$, while in Figure 2, $N = 9$ and $M = 4$. Then we can write $y_t = Sb_t$, where y_t is a vector of all the level nodes at time t and b_t is the vector of all the bottom level nodes at time t .




Figure 2: An example of a two level grouped structure.

For the example shown in Figure 2, the equation can be written as follows:

$$\begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \\ y_{D,t} \\ y_{AC,t} \\ y_{AD,t} \\ y_{BC,t} \\ y_{BD,t} \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} y_{AC,t} \\ y_{AD,t} \\ y_{BC,t} \\ y_{BD,t} \end{pmatrix}.$$

1.2 Forecasting hierarchical time series

If we just forecast each series individually, we are ignoring the hierarchical or grouping structure, and the forecasts will not be “coherent” (they will not add up appropriately).

There are several available methods that consider the hierarchical structure information when forecasting time series. These include the top-down (Gross & Sohl 1990; Fliedner 2001), bottom-up (Kahn 1998), middle-out and optimal combination (Hyndman et al. 2011) approaches. In the top-down approach, we first forecast the total series and then disaggregate the forecast to form lower level series forecasts based on a set of historical and forecasted proportions (for details see Athanasopoulos, Ahmed & Hyndman 2009). In the bottom-up approach, the forecasts in each level of the hierarchy can be computed by aggregating the bottom level series forecasts. However, we may not get good upper-level forecasts because the most disaggregated series can

be noisy and so their forecasts are often inaccurate. In the middle-out approach, the process can be started from one of the middle levels and other forecasts can be computed using aggregation for upper levels and disaggregation for lower levels. Finally, optimal combination uses all the n forecasts for all of the series in the entire structure, and then uses an optimization process to reconcile the resulting forecasts. The advantage of the optimal combination method, compared with the other methods, is that it considers all information in the hierarchy, including any correlations among the series.

In the optimal combination method, reconciled forecasts can be computed using the following equation known as weighted least squares (WLS) (Wickramasuriya, Athanasopoulos & Hyndman 2019)

$$\tilde{\mathbf{y}}_h = \mathbf{S}(\mathbf{S}'\mathbf{W}_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\mathbf{W}_h^{-1}\hat{\mathbf{y}}_h, \quad (1)$$

where $\hat{\mathbf{y}}_h$ represents a vector of h -step-ahead base forecasts for all levels of the hierarchy, and \mathbf{W}_h is the covariance matrix of forecast errors for the h -step-ahead base forecasts.

Several possible simple methods for estimating \mathbf{W}_h are available. Wickramasuriya, Athanasopoulos & Hyndman (2019) discuss a simple approximation whereby $\mathbf{W}_h = k_h\mathbf{\Lambda}$ with k_h being a positive constant, $\mathbf{\Lambda} = \text{diag}(\mathbf{S}\mathbf{1})$, and $\mathbf{1}$ being a column of 1s. Note that $\mathbf{\Lambda}$ simply contains the row sums of the summing matrix \mathbf{S} , and that k_h will cancel out in (1). Thus

$$\tilde{\mathbf{y}}_h = \mathbf{S}(\mathbf{S}'\mathbf{\Lambda}^{-1}\mathbf{S})^{-1}\mathbf{S}'\mathbf{\Lambda}^{-1}\hat{\mathbf{y}}_h. \quad (2)$$

The most computationally challenging part of the optimal combination method is to produce all the base forecasts that make up $\hat{\mathbf{y}}_h$. In many applications, there may be thousands or even millions of individual series, and each of them must be forecast independently. The most popular time series forecasting methods such as ETS and ARIMA models (Hyndman & Athanasopoulos 2018) involve non-linear optimization routines to estimate the parameters via maximum likelihood estimation. Usually, multiple models are fitted for each series, and the best is selected by minimizing Akaike's Information Criterion (Akaike 1998). This computational challenges increases with the number of lower level series as well as in the number of aggregations of interest.

We therefore propose a new approach to compute the base forecasts that is both computationally fast while maintaining an acceptable forecasting accuracy level.

2 Proposed approach: Linear model

Our proposed approach is based on using linear regression models for computing base forecasts. Suppose we have a linear model that we use for forecasting, and we wish to apply it to N different series which have some aggregation constraints. We have observations $y_{t,i}$ from times $t = 1, \dots, T$ and series $i = 1, \dots, N$. Then

$$y_{t,i} = \beta_i' x_{t,i} + \varepsilon_{t,i}$$

where $x_{t,i} = \{1, x_{t,i,1}, \dots, x_{t,i,p}\}$ is a $(p+1)$ -vector of regression variables. This equation for all the observations in matrix form can be written as follows:

$$\begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{pmatrix} = \begin{pmatrix} X_1 & 0 & 0 & \dots & 0 \\ 0 & X_2 & 0 & \dots & 0 \\ 0 & 0 & X_3 & \ddots & \vdots \\ \vdots & \vdots & \ddots & \ddots & 0 \\ 0 & 0 & \dots & 0 & X_N \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \\ \vdots \\ \beta_N \end{pmatrix} + \begin{pmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \varepsilon_3 \\ \vdots \\ \varepsilon_N \end{pmatrix}, \quad (3)$$

where $y_i = \{y_{1,i}, y_{2,i}, \dots, y_{T,i}\}$ is a T -vector, $\beta_i = \{\beta_{0,i}, \beta_{1,i}, \beta_{2,i}, \dots, \beta_{p,i}\}$ is a $(p+1)$ -vector, $\varepsilon_i = \{\varepsilon_{1,i}, \varepsilon_{2,i}, \dots, \varepsilon_{T,i}\}$ is a T -vector and X_i is the $T \times (p+1)$ -matrix

$$X_i = \begin{pmatrix} 1 & x_{1,i,1} & x_{1,i,2} & \dots & x_{1,i,p} \\ 1 & x_{2,i,1} & x_{2,i,2} & \dots & x_{2,i,p} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{T,i,1} & x_{T,i,2} & \dots & x_{T,i,p} \end{pmatrix}.$$

Equation (3) can be written as $Y = XB + E$, with parameter estimates given by $\hat{B} = (X'X)^{-1}X'Y$.

Then the base forecasts are obtained using

$$\hat{y}_{t+h} = X_{t+h}^* \hat{B}, \quad (4)$$

where \hat{y}_{t+h} is an N -vector of forecasts, \hat{B} comprises N stacked $(p+1)$ -vectors of estimated coefficients, and X_{t+h}^* is the $N \times N(p+1)$ matrix

$$\mathbf{X}_{t+h}^* = \begin{pmatrix} \mathbf{x}'_{t+h,1} & 0 & 0 & \dots & 0 \\ 0 & \mathbf{x}'_{t+h,2} & 0 & \dots & 0 \\ 0 & 0 & \mathbf{x}'_{t+h,3} & \ddots & \vdots \\ \vdots & \vdots & \ddots & \ddots & 0 \\ 0 & 0 & \dots & 0 & \mathbf{x}'_{t+h,N} \end{pmatrix}.$$

Note that we use \mathbf{X}_t^* to distinguish this matrix, which combines $\mathbf{x}_{t,i}$ across all series for one time from \mathbf{X}_i which combines $\mathbf{x}_{t,i}$ across all time for one series.

Finally, we can combine the two linear equations for computing base forecasts and reconciled forecasts (Equations (2) and (4)) to obtain the reconciled forecasts with a single equation:

$$\tilde{\mathbf{y}}_{t+h} = \mathbf{S}(\mathbf{S}'\mathbf{\Lambda}\mathbf{S})^{-1}\mathbf{S}'\mathbf{\Lambda}(\mathbf{X}_{t+h}^*\hat{\mathbf{B}}) = \mathbf{S}(\mathbf{S}'\mathbf{\Lambda}\mathbf{S})^{-1}\mathbf{S}'\mathbf{\Lambda}\mathbf{X}_{t+h}^*(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}. \quad (5)$$

2.1 Simplified formulation for a fixed set of predictors (X)

If we have the same set of predictor variables, \mathbf{X} , for all the series, we can write Equations (3) to (5) more easily using multivariate regression equations, and we can obtain all the reconciled forecasts for all the series in one equation. In that case, Equation (3) can be rearranged as follows:

$$\begin{pmatrix} y_{11} & \dots & y_{1N} \\ y_{21} & \dots & y_{2N} \\ \vdots & & \vdots \\ y_{T1} & \dots & y_{TN} \end{pmatrix} = \begin{pmatrix} 1 & X_{11} & \dots & X_{1p} \\ 1 & X_{21} & \dots & X_{2p} \\ \vdots & \vdots & & \vdots \\ 1 & X_{T1} & \dots & X_{Tp} \end{pmatrix} \begin{pmatrix} \beta_{01} & \dots & \beta_{0N} \\ \beta_{11} & \dots & \beta_{1N} \\ \vdots & & \vdots \\ \beta_{p1} & \dots & \beta_{pN} \end{pmatrix} + \begin{pmatrix} \varepsilon_{11} & \dots & \varepsilon_{1N} \\ \varepsilon_{21} & \dots & \varepsilon_{2N} \\ \vdots & & \vdots \\ \varepsilon_{T1} & \dots & \varepsilon_{TN} \end{pmatrix}, \quad (6)$$

where \mathbf{Y} , \mathbf{X} , \mathbf{B} and \mathbf{E} are now matrices of size $T \times N$, $T \times (p+1)$, $(p+1) \times N$ and $T \times N$, respectively. Equations (4) to (5) can be written accordingly using Equation (6) and here $\mathbf{X}_{t+h,i}^* = \mathbf{X}_{t+h,i}'$, where \mathbf{X}_{t+h}^* is an $h \times (p+1)$ matrix.

2.2 OLS predictors

As an example of the \mathbf{X}_t matrix in Equation (3), we can refer to the set of predictors proposed in Ashouri, Shmueli & Sin (2018) for modeling trend, seasonality and autocorrelation by using lagged values (y_{t-1}, y_{t-2}, \dots), trend variables and seasonal dummy variables:

$$y_t = \alpha_0 + \alpha_1 t + \beta_1 s_{1,t} + \dots + \beta_{m-1} s_{m-1,t} + \gamma_1 y_{t-1} + \dots + \gamma_p y_{t-p} + \delta z_t + \varepsilon_t. \quad (7)$$

Here, $s_{j,t}$ is a dummy variable taking value 1 if time t is in season j ($j = 1, 2, \dots, m$), y_{t-k} is the k th lagged value for y_t and z_t is some external information at time t . The seasonal period m

depends on the problem; for instance, if we have daily data with day-of-week seasonality, then $m = 7$.

Because of using lags and external series as predictors in Equation (7), we do not have same set of predictors for all the series, y_t . However, if we just use trend and seasonality dummies as the predictors, then the simpler equations given in Section 2 can be written using multivariate regression models.

While OLS is popular in practice for forecasting time series, it is often frowned upon due to its independence assumption. This can cause issues for parametric inference but is less of a problem for forecasting. In fact it often performs sufficiently well for forecasting as can be seen by its popular use in practice. Further, the use of autoregressive terms in the above model should model most of the autocorrelation in the data.

2.3 Computational considerations

There are two ways for computing the above forecasts. First, we could create the matrices Y , X and E , and then directly use the above equations (taking advantage of sparse matrix routines) to obtain the forecasts. Alternatively, to avoid numerical issues due to inverting a large matrix, we could use separate regression models to compute the coefficients for each linear model individually. Although the matrix, $X'X$, which we need to invert is sparse and block diagonal, it is still faster and more accurate to use the second approach involving separate regression models. In Appendix A we give the forecast results comparisons for Australian domestic tourism example using these two computation approaches.

2.4 Prediction intervals

For obtaining prediction intervals, we need to compute the variance of reconciled forecasts as follows (Wickramasuriya, Athanasopoulos & Hyndman 2019):

$$\text{var}(\tilde{y}_{t+h}) = S(S'\Lambda S)^{-1}S'\Lambda\Sigma_{t+h}((S'\Lambda S)^{-1}S'\Lambda)'S' = SP\Sigma_{t+h}P'S'$$

where Σ_{t+h} represents the variance of the base forecasts and $P = (S'\Lambda S)^{-1}S'\Lambda$.

3 Applications

In this section we illustrate our approach using two examples: forecasting monthly Australian domestic tourism and forecasting daily Wikipedia pageviews. We compare the forecasting accuracy of ETS, ARIMA and the proposed linear OLS forecasting model, with and without the reconciliation step. In these applications, we used the weighted reconciliation approach

from Equation (2). For comparing these methods, we use the average of Root Mean Square Errors (RMSEs) across all series and also display box plots for forecast errors along with the raw forecast errors.

The two datasets differ in terms of structure, size and behavior. The tourism data contains 304 series with both hierarchical and grouped structure, while the Wikipedia pageviews dataset contains 913 series with grouped structure. The tourism dataset has strong seasonality while the Wikipedia data are noiser.

We apply two methods for generating forecasts, which differ in how they handle unobserved lagged values as inputs. The first approach is *ex post* in that it uses actual values, even when they are future to the forecast origin. These values are known to us because they are in the test set. We call these *rolling origin* forecasts. In the second *ex ante* method, we replace lagged values of y by their forecasts if they occur at periods after the forecast origin. We call these *fixed origin* forecasts.

3.1 Australian domestic tourism

This dataset has 19 years of monthly visitor nights in Australia by Australian tourists, a measure used as an indicator of tourism activity (Wickramasuriya, Athanasopoulos & Hyndman 2019). The data were collected by computer-assisted telephone interviews with 120,000 Australians aged 15 and over (Tourism Research Australia 2005). The dataset includes 304 time series each of length 228 observations. The hierarchy and grouping structure for this dataset is made using geographic and purpose of travel information.

Table 1: *Australia geographic hierarchical structure.*

Series	Name	Label	Series	Name	Label
Total			Region		
1	Australia	Total	55	Lakes	BCA
State			56	Gippsland	BCB
2	NSW	A	57	Phillip Island	BCC
3	VIC	B	58	General Murray	BDA
4	QLD	C	59	Goulburn	BDB
5	SA	D	60	High Country	BDC
6	WA	E	61	Melbourne East	BDD
7	TAS	F	62	Upper Yarra	BDE
8	NT	G	63	Murray East	BDF
Zone			64	Wimmera+Mallee	BEA
9	Metro NSW	AA	65	Western Grampians	BEB
10	Nth Coast NSW	AB	66	Bendigo Loddon	BEC

	11	Sth Coast NSW	AC	67	Macedon	BED
	12	Sth NSW	AD	68	Spa Country	BEE
	13	Nth NSW	AE	69	Ballarat	BEF
	14	ACT	AF	70	Central Highlands	BEG
	15	Metro VIC	BA	71	Gold Coast	CAA
	16	West Coast VIC	BB	72	Brisbane	CAB
	17	East Coast VIC	BC	73	Sunshine Coast	CAC
	18	Nth East VIC	BD	74	Central Queensland	CBA
	19	Nth West VIC	BE	75	Bundaberg	CBB
	20	Metro QLD	CA	76	Fraser Coast	CBC
	21	Central Coast QLD	CB	77	Mackay	CBD
	22	Nth Coast QLD	CC	78	Whitsundays	CCA
	23	Inland QLD	CD	79	Northern	CCB
	24	Metro SA	DA	80	Tropical North Queensland	CCC
	25	Sth Coast SA	DB	81	Darling Downs	CDA
	26	Inland SA	DC	82	Outback	CDB
	27	West Coast SA	DD	83	Adelaide	DAA
	28	West Coast WA	EA	84	Barossa	DAB
	29	Nth WA	EB	85	Adelaide Hills	DAC
	30	Sth WA	EC	86	Limestone Coast	DBA
	31	Sth TAS	FA	87	Fleurieu Peninsula	DBB
	32	Nth East TAS	FB	88	Kangaroo Island	DBC
	33	Nth West TAS	FC	89	Murraylands	DCA
	34	Nth Coast NT	GA	90	Riverland	DCB
	35	Central NT	GB	91	Clare Valley	DCC
Region				92	Flinders Range and Outback	DCD
	36	Sydney	AAA	93	Eyre Peninsula	DDA
	37	Central Coast	AAB	94	Yorke Peninsula	DDB
	38	Hunter	ABA	95	Australia’s Coral Coast	EAA
	39	North Coast NSW	ABB	96	Experience Perth	EAB
	40	Northern Rivers Tropical NSW	ABC	97	Australia’s SouthWest	EAC
	41	South Coast	ACA	98	Australia’s North West	EBA
	42	Snowy Mountains	ADA	99	Australia’s Golden Outback	ECA
	43	Capital Country	ADB	100	Hobart and the South	FAA
	44	The Murray	ADC	101	East Coast	FBA
	45	Riverina	ADD	102	Launceston, Tamar and the North	FBB
	46	Central NSW	AEA	103	North West	FCA
	47	New England North West	AEB	104	Wilderness West	FCB
	48	Outback NSW	AEC	105	Darwin	GAA
	49	Blue Mountains	AED	106	Kakadu Arnhem	GAB
	50	Canberra	AFA	107	Katherine Daly	GAC
	51	Melbourne	BAA	108	Barkly	GBA

52	Peninsula	BAB	109	Lasseter	GBB
53	Geelong	BAC	110	Alice Springs	GBC
54	Western	BBA	111	MacDonnell	GBD

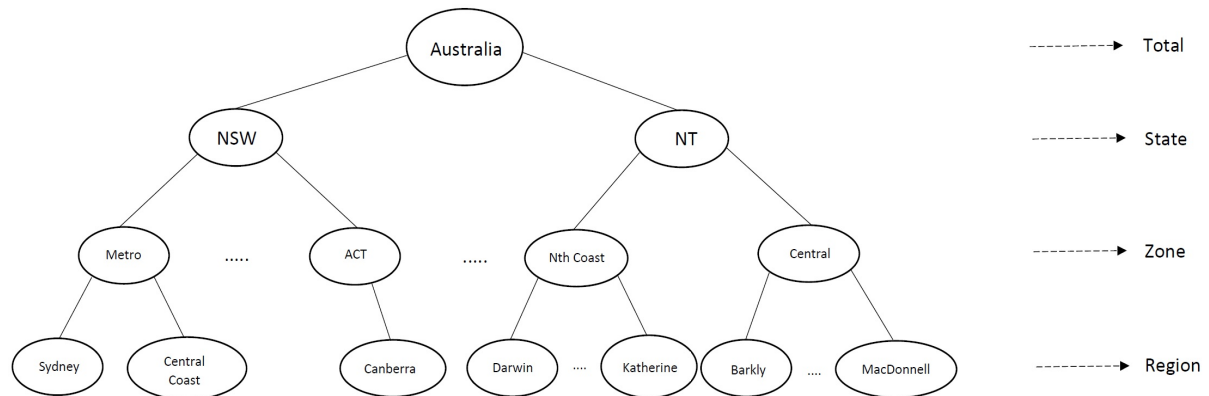


Figure 3: Australian geographic hierarchical structure.




Figure 4: Australia tourism region map - colors represent states.

In this dataset we have three levels of geographic divisions in Australia. In the first level, Australia is divided into seven ‘States’ including New South Wales (NSW), Victoria (VIC),

Queensland (QLD), South Australia (SA), Western Australia (WA), Tasmania (TAS) and Northern Territory (NT). In the second and third levels, it is divided into 27 ‘Zones’ and 76 ‘Regions’ (for details about Australia geographic divisions see Figure 3 and Table 1 and also Figure 4 which shows Australia map divided by tourism region and colored by states¹).

For ‘Purpose’ we have four groups: Holiday (Hol), Visiting friends and relatives (Vis), Business (Bus) and Other (Oth). Based on the geographic hierarchy and purpose grouping, we end up with 8 hierarchical levels with 555 series in total:

- Level 0 = Total for Australia
- Level 1 = State totals
- Level 2 = Zone totals
- Level 3 = Region totals
- Level 4 = Purpose totals
- Level 5 = State * Purpose totals
- Level 6 = Zone* Purpose totals
- Level 7 = bottom level series

Table 2: Number of Australian domestic tourism series in each level of the hierarchy and group structure.

Geographic division	# series (geographic division)	# series (purpose of travel)	Total
Australia	1	4	5
State	7	28	35
Zone	27	108	135
Region	76	304	380
Total	111	444	555

We report the forecast results for all these hierarchical levels, as well as the average RMSE across all the levels of the hierarchy. We used different predictors in the OLS predictor matrix for the rolling and fixed origin approaches. For the rolling origin model, we included a linear trend, 11 dummy variables, and 12 lags. For the fixed origin model, we included a quadratic trend, 11 dummy variables, and lags 1 and 12. This is intended to capture the monthly seasonality. In addition, before running the model, we partition the data into training and test sets, with the last 24 months (2 years) as our test set, and the rest as our training set.

In Tables 3 and 4, we show the average RMSE for the test set forecasts. Methods include ETS, ARIMA and our proposed OLS forecasting model. In Table 3 we forecast 24 1-step-ahead

¹For more information about the map: https://www.tra.gov.au/tra/2016/Tourism_Region_Profiles/Region_profiles/index.html

Table 3: Mean(RMSE) on 2 year test set for ETS, ARIMA and OLS with and without reconciliation - Rolling origin - Tourism dataset

Level	Unreconciled			Reconciled		
	ETS	ARIMA	OLS	ETS	ARIMA	OLS
Total	1516.4	1445.5	1415.1	1517.2	1517.2	1414.7
State	511.4	493.1	510.8	499.9	499.9	491.6
Zone	214.8	219.0	224.5	209.6	209.6	213.6
Region	122.9	125.1	124.0	119.4	119.4	120.4
Purpose	676.0	709.2	694.5	674.2	674.2	681.9
State*Purpose	213.1	220.1	216.1	212.7	212.7	213.3
Zone*Purpose	97.5	102.4	101.0	96.8	96.8	99.0
Bottom level	56.2	58.2	58.2	56.2	56.2	57.4

Table 4: Mean(RMSE) on 2 year test set for ETS, ARIMA and OLS with and without reconciliation - Fixed origin - Tourism dataset

Level	Unreconciled			Reconciled		
	ETS	ARIMA	OLS	ETS	ARIMA	OLS
Total	2238.6	3554.0	2528.9	2232.8	3460.3	2540.1
State	593.6	570.1	596.5	555.7	658.5	579.0
Zone	239.5	229.6	243.3	235.3	249.8	235.4
Region	132.6	129.4	127.1	127.6	132.4	123.8
Purpose	766.8	824.0	875.5	801.7	1019.3	857.2
State*Purpose	226.7	241.2	236.7	224.5	245.6	229.1
Zone*Purpose	103.0	105.4	104.9	102.4	105.8	102.9
Bottom level	59.1	58.8	58.6	58.8	59.3	57.9

forecasts for the test set, rolling forward month by month. In Table 4 we generate fixed origin forecasts up to 24 steps ahead. In these tables we show forecasts with and without reconciliation.

These results show that our proposed OLS forecasting model produces forecast accuracy similar to ETS and ARIMA, which are computationally heavy for many time series. Also they show the usefulness of the reconciliation in decreasing the average RMSE in all three methods. Except for the total series, reconciliation has helped in forecasting all the hierarchical levels. Also, because the higher level series have higher counts, the errors are larger in magnitude (see Figures 5 and 6). To better compare errors across all the hierarchy levels we scaled the errors (See Appendix B).

In Figures 5 and 6 we display the error box plots for both reconciled and unreconciled forecasts using all three methods, for the rolling origin and fixed origin forecasts. In these figures we can visualize the error distributions across all the models, as well as the usefulness of the reconciliation step in improving the forecasts. In particular, we see that (as expected) by

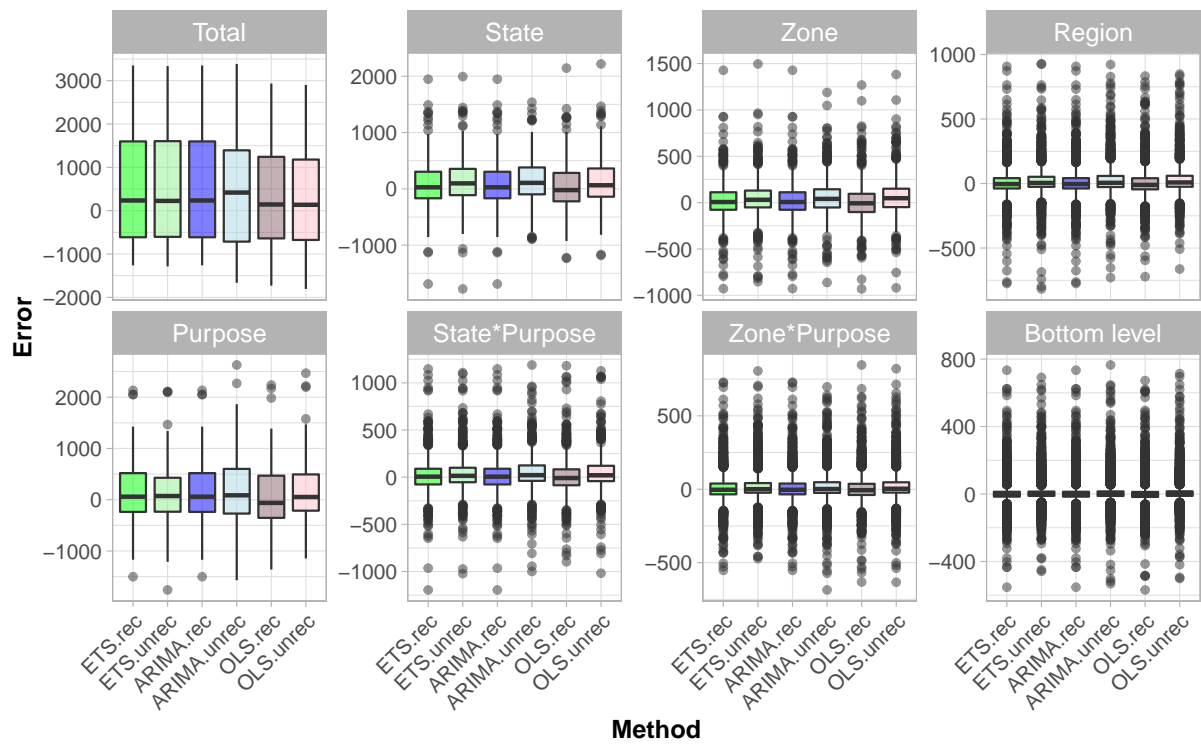


Figure 5: Box plots of rolling origin forecast errors from reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for tourism demand.

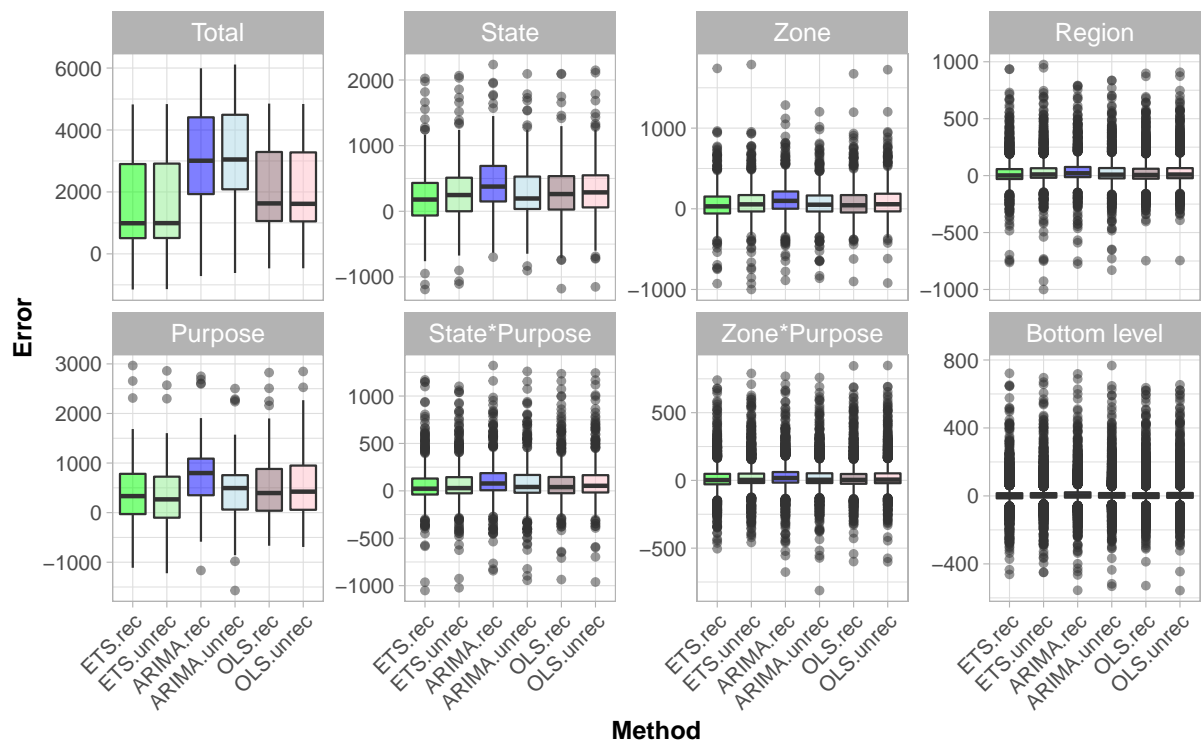


Figure 6: Box plots of fixed origin forecast errors for reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for tourism demand.

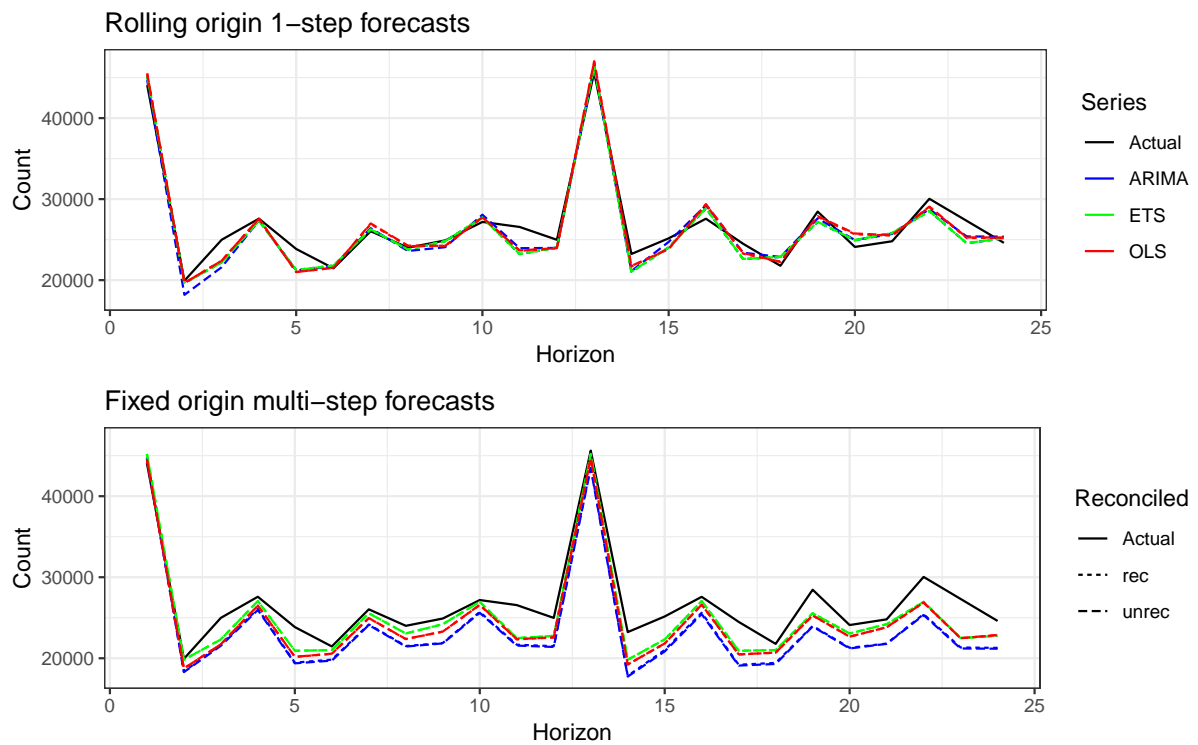


Figure 7: The actual test set for the 'Total series' compared to the forecasts from reconciled and unreconciled ETS, ARIMA and OLS methods for rolling and fixed origin tourism demand.

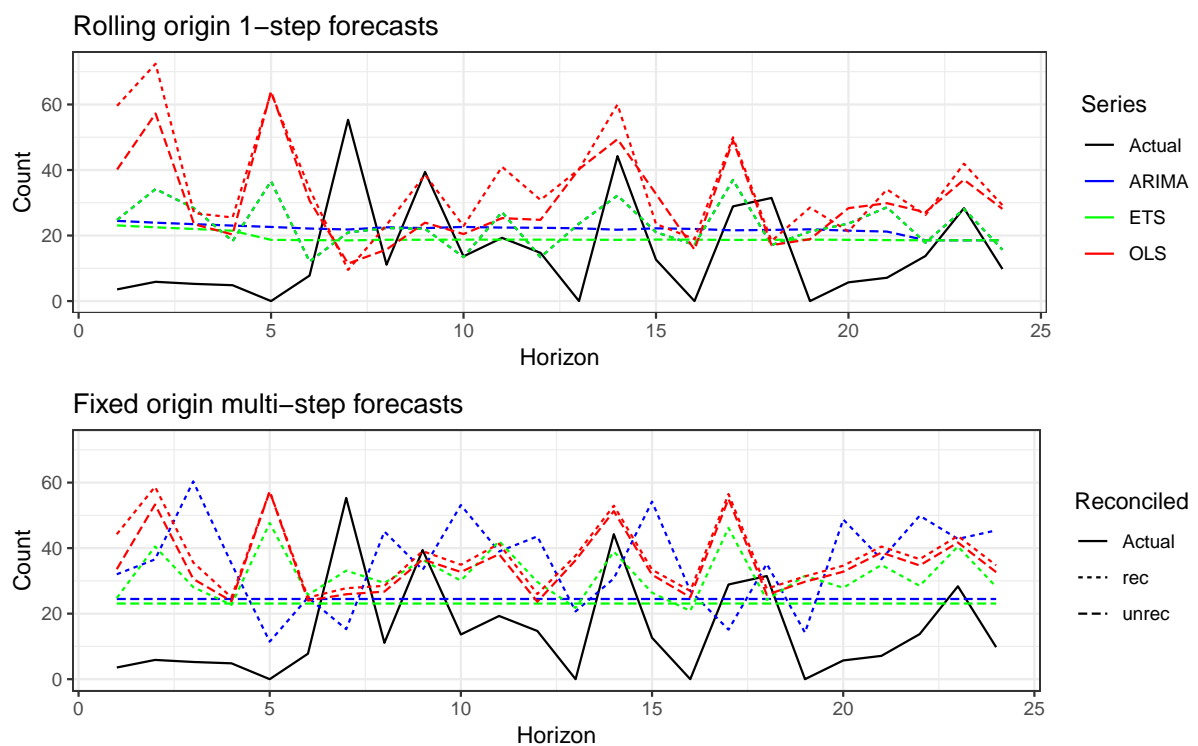


Figure 8: The actual test set for the 'BACBus' bottom level series compared to the forecasts from reconciled and unreconciled ETS, ARIMA and OLS methods for rolling and fixed origin tourism demand.

applying rolling origin 1-step-ahead forecasts, the error densities are closer and more tightly distributed around zero than the fixed origin multi-step-ahead forecasts.

Figures 7 and 8 show the rolling and fixed origin forecast results for the total series and one of the bottom level series, BACBus (Geelong - Business). In these plots we have both reconciled (solid lines) and unreconciled (dashed lines) forecasts and we see that the reconciliation step improves the forecasts in this series. We also see that the OLS model forecast accuracy is similar to the other two methods.

Table 5 compares the computation time of the three methods for rolling and fixed origin forecasting. We see that the OLS forecasting model is much faster compared to the other methods. Also, since reconciliation is a linear process, in all methods it is very fast and does not affect computation time significantly.

Table 5: *Computation time (seconds) for ETS, ARIMA and OLS with and without reconciliation - Rolling and fixed origin forecasts on a 24 month test set - Tourism dataset*

	Rolling origin		Fixed origin	
	Unreconciled	Reconciled	Unreconciled	Reconciled
ETS	10924.57	10924.60	407.10	407.15
ARIMA	31146.38	31146.52	1116.15	1116.19
OLS	48.40	48.31	17.42	17.80

Since we are using a linear model, we can easily include exogenous variables which can often be helpful in improving forecast accuracy. In this application, we tried including an “Easter” dummy variable indicating the timing of Easter, but its affect on forecast accuracy was minimal, so it was omitted in the model reported here.

3.2 Wikipedia pageviews: Grouped structure

The second dataset comprises one year of daily data (2016-06-01 to 2017-06-29) on Wikipedia pageviews for the most popular social networks articles (Ashouri, Shmueli & Sin 2018). This dataset is noisier than the Australian monthly tourism data, making forecasting more challenging. The data has a grouped structure with the following attributes: ‘Agent’: Spider, User, ‘Access’: Desktop, Mobile app, Mobile web, ‘Language’: en (English), de (German), es (Spanish), zh (Chinese) and ‘Purpose’: Blogging related, Business, Gaming, General purpose, Life style, Photo sharing, Reunion, Travel, Video (see Table 6).

Table 6: *Social networking Wikipedia article grouping structure*

Grouping	Series	Grouping	Series
Total		Language	
Access	1. Social Network	Purpose	10. zh (Chinese)
Agent	2. Desktop		11. Blogging related
	3. Mobile app		12. Business
	4. Mobile web		13. Gaming
	5. Spider		14. General purpose
	6. User		15. Life style
Language	7. en (English)		16. Photo sharing
	8. de (German)		17. Reunion
	9. es (Spanish)		18. Travel
			19. Video

The final dataset includes 913 time series, each with length 394. The applied group structure and different levels include²:

- Level 0 = Total
- Level 1 = Access
- Level 2 = Agent
- Level 3 = Language
- Level 4 = Purpose
- Level 5 = Access * Agent
- Level 6 = Access * Language
- Level 7 = Access * Purpose
- Level 8 = Agent * Language
- Level 9 = Agent * Purpose
- Level 10 = Language * Purpose
- Level 11 = bottom level series

For this daily dataset, in the OLS forecasting model we include in the predictor matrix a quadratic trend, 6 seasonal dummies and all 7 lags for rolling, and for fixed origin model we use a quadratic trend, 6 seasonal dummies and lags 1 and 7. We partitioned the data into two parts training and test sets. We used the last 28 days for our test set and the rest for the training set. In this example, the results in tables and figures are represented for single groups although we applied all the above levels in the group structure for reconciliation.

²There are 4 more levels aggregates for Agent*Access*Language, Agent*Access*Purpose, Agent*Language*Purpose, and Access*Language*Purpose. Including these 4 additional levels might slightly improve the results but because of simplicity we applied mentioned 12 levels in the text.

Tables 7 and 8 show the RMSE results. Although these time series are noisier, we still get acceptable results for the OLS forecasting model compared with ETS and ARIMA. In this case, we get similar results with and without the reconciliation step.

Table 7: *Mean(RMSE) for ETS, ARIMA and OLS with and without reconciliation - Rolling origin - Wikipedia dataset*

Level	Unreconciled			Reconciled		
	ETS	ARIMA	OLS	ETS	ARIMA	OLS
Total	10524.6	15648.5	17066.5	9782.5	15123.5	16695.1
Access	6757.6	5222.6	4395.9	6619.5	6603.5	4219.0
Agent	8518.3	9894.2	13043.7	8015.8	10295.6	12902.8
Language	3671.8	6662.7	5677.0	3996.6	6616.4	5623.9
Purpose	4370.0	3356.9	3295.8	4575.6	3461.1	3264.7
Bottom level	374.7	213.6	231.3	405.3	219.1	227.3

Table 8: *Mean(RMSE) for ETS, ARIMA and OLS with and without reconciliation - Fixed origin - Wikipedia dataset*

Level	Unreconciled			Reconciled		
	ETS	ARIMA	OLS	ETS	ARIMA	OLS
Total	20374.7	23185.6	27322.8	21302.4	23585.7	27168.6
Access	4505.7	10853.0	6930.1	5115.3	11736.6	7101.6
Agent	20238.3	16704.7	21059.8	18637.0	16553.3	20945.8
Language	6368.1	10611.3	8696.5	6755.0	10908.9	8900.7
Purpose	4848.5	8969.1	4574.9	4618.8	8411.0	4557.4
Bottom level	429.8	362.0	328.7	443.2	367.0	328.4

Figures 9 and 10 display the forecast error box plot. These plots are for rolling and fixed origin forecasts over 28 days in each level of grouping. Further, we can see that the error distribution is almost similar in all levels across the different methods. The only exception is the Total series, where ETS performs significantly better than ARIMA and OLS. We also note that the reconciliation is less effective. As in the tourism example, in higher levels, series have higher counts and therefore their error magnitudes are larger.

In Figures 11 and 12, we display results for the total and one of the bottom level series, ‘desktopusenPho’ (desktop-user-english-photo sharing). The plot shows rolling and fixed origin forecast results over the 28 day test set for ETS, ARIMA and OLS, with (solid lines) and without (dashed lines) applying reconciliation. We see that the OLS forecasting model performs close to the other two methods, and reconciliation improves the forecasts.

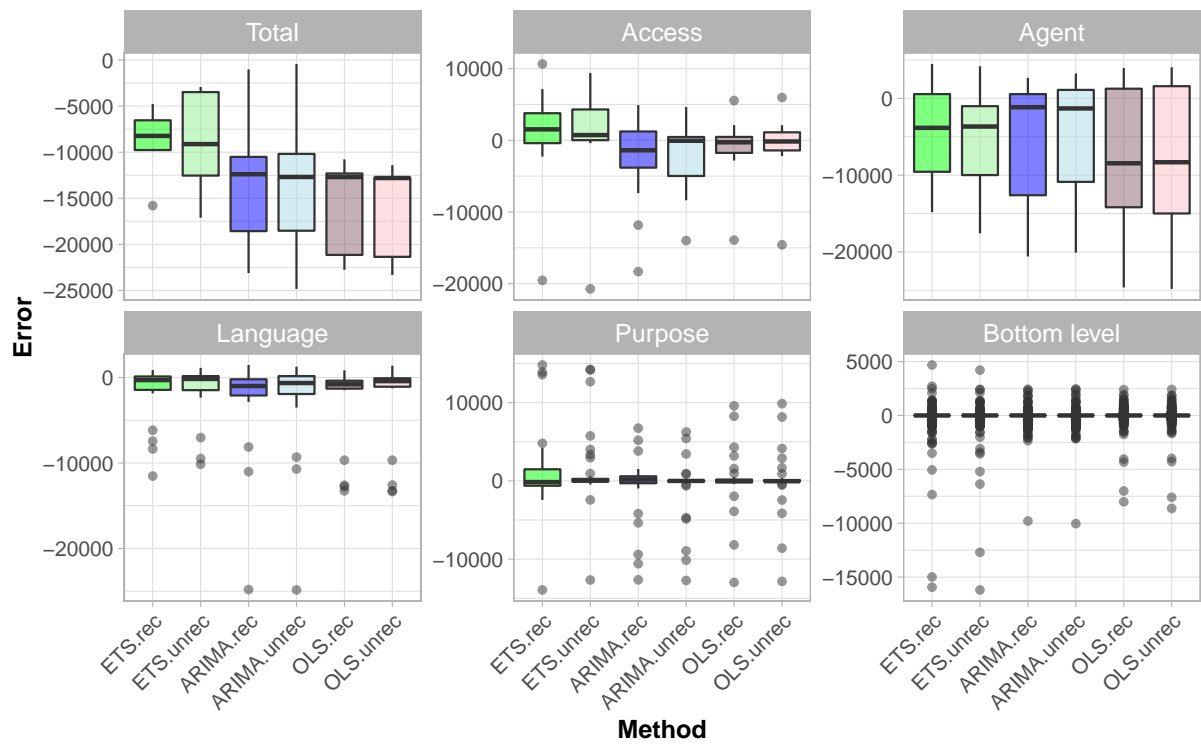


Figure 9: Box plots of forecast errors for reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for rolling origin forecasts of Wikipedia pageviews.

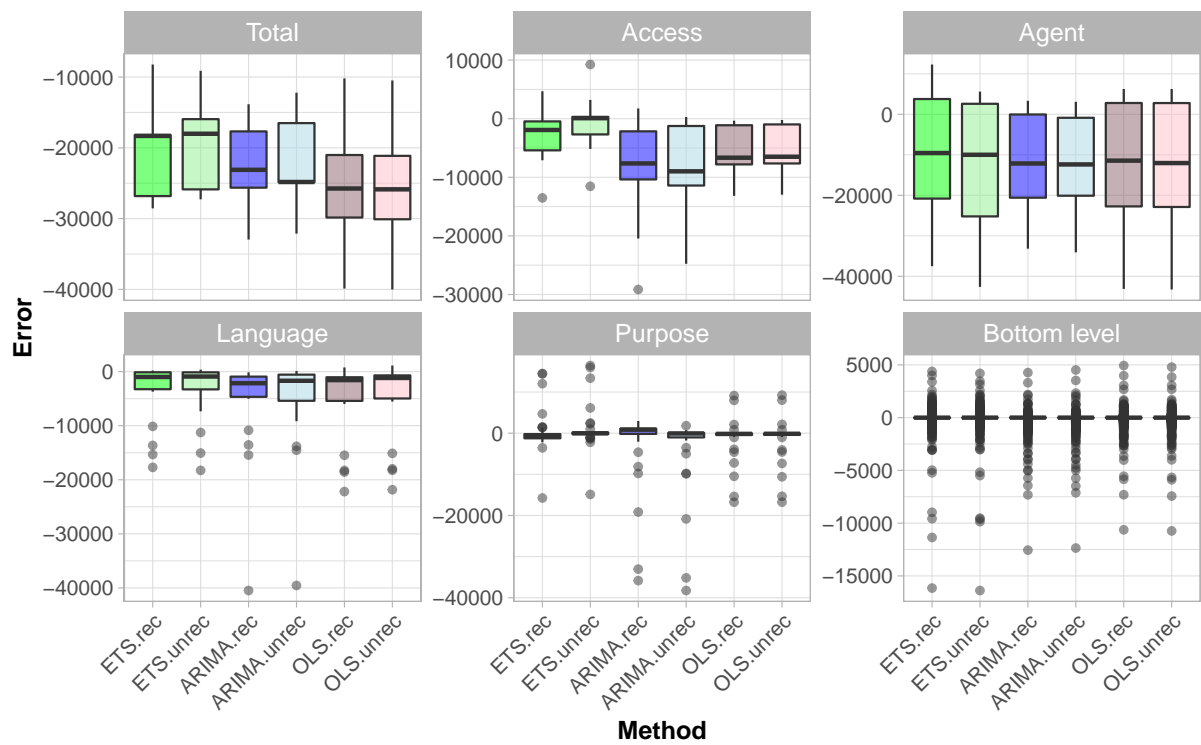


Figure 10: Box plots of forecast errors for reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for fixed origin forecasts of Wikipedia pageviews.

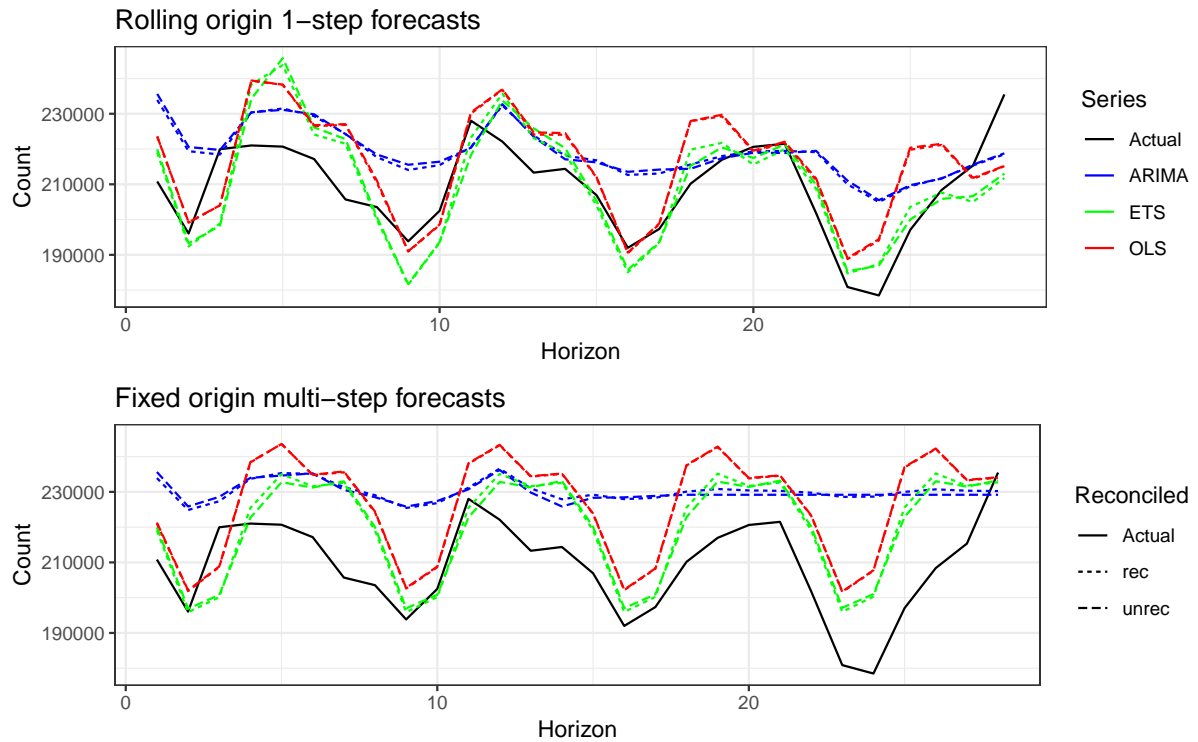


Figure 11: The actual test set for the 'Total' series compared to the forecasts from reconciled and unreconciled ETS, ARIMA and OLS methods for rolling and fixed origin forecasts of Wikipedia pageviews.




Figure 12: The actual test set for the 'desktopusenPho02' bottom level series compared to the forecasts from reconciled and unreconciled ETS, ARIMA and OLS methods for rolling and fixed origin forecasts of Wikipedia pageviews.

Table 9 presents the computation times for all three methods. ETS and ARIMA are clearly much more computationally heavy compared with OLS. As in the Australian tourism dataset, running reconciliation does not have much effect on computation time.

Table 9: *Computation time (seconds) for ETS, ARIMA and OLS with and without reconciliation - Rolling and fixed origin forecasts - Wikipedia dataset*

	Computation time (secs)			
	Rolling origin		Fixed origin	
	Unreconciled	Reconciled	Unreconciled	Reconciled
ETS	20613.08	20613.14	871.55	871.58
ARIMA	29419.36	29419.39	1169.52	1169.56
OLS	116.27	116.31	61.33	61.38

4 Conclusion

We have proposed a linear model approach to fast forecasting of hierarchical or grouped time series, with accuracy that nearly matches that of forecast methods such as ETS and ARIMA. This is especially useful in large collections of time series, as is typical in hierarchical and grouped structures. Although ETS and ARIMA are advantageous in terms of forecasting power and accuracy, they can be computationally heavy when facing large collections of time series in the hierarchy. An important feature of our model is its ability to easily include external information such as holiday dummies or other external series. We also note that OLS has the additional practical advantage of handling missing data while ETS requires imputation.

Another advantage of our approach is that it can be computed in a single matrix equation (5). This makes it extremely fast and easy to implement, and enables standard results to be derived with minimal effort (e.g., prediction intervals).

Pennings & Dalen (2017) proposed another approach for forecasting hierarchical time series using state space models. Although their approach is flexible in handling outliers, missing data and external features, it is less flexible to different kinds of datasets and it is computationally much more complex.

Acknowledgements

The first and third authors of this research were partially funded by Ministry of Science and Technology (MOST), Taiwan [Grant 106-2420-H-007-019].

Appendix A

We compare the forecasts obtained using separate regression models versus matrix computation for the Australian domestic tourism example. The two methods give numerically equivalent results, but Table 10 shows the matrix approach is computationally heavier (even with sparse matrix algebra).

Table 10: *Computation time (seconds) for OLS using the matrix approach and separate regression models, with and without reconciliation, on a rolling and fixed origin for 24 steps ahead, using the tourism dataset.*

	Rolling origin		Fixed origin	
	Unreconciled	Reconciled	Unreconciled	Reconciled
Matrix approach	202.06	209.84	87.73	105.69
Separate models	48.40	48.31	16.66	16.85

Appendix B

We provide boxplots of the scaled forecasted errors for the tourism example. These plots are displayed for both rolling forward and multiple-step-ahead forecasts.




Figure 13: Box plots of scaled forecast errors from reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for rolling origin 1-step-ahead tourism demand.




Figure 14: Box plots of scaled forecast errors from reconciled and unreconciled ETS, ARIMA and OLS methods at each hierarchical level for fixed origin multi-step-ahead tourism demand.

References

- Akaike, H (1998). “Information theory and an extension of the maximum likelihood principle”. In: *Selected Papers of Hirotugu Akaike*. Springer Series in Statistics (Perspectives in Statistics). Springer, pp.199–213.
- Ashouri, M, G Shmueli & CY Sin (2018). Clustering time series by domain-relevant features using model-based trees. *Proceedings of the 2018 Data Science, Statistics & Visualization (DSSV)*.
- Athanasopoulos, G, RA Ahmed & RJ Hyndman (2009). Hierarchical forecasts for Australian domestic tourism. *International Journal of Forecasting* **25**(1), 146–166.
- Fliedner, G (2001). Hierarchical forecasting: issues and use guidelines. *Industrial Management & Data Systems* **101**(1), 5–12.
- Gross, CW & JE Sohl (1990). Disaggregation methods to expedite product line forecasting. *Journal of Forecasting* **9**(3), 233–254.
- Hyndman, RJ, RA Ahmed, G Athanasopoulos & HL Shang (2011). Optimal combination forecasts for hierarchical time series. *Computational Statistics & Data Analysis* **55**(9), 2579–2589.
- Hyndman, RJ & G Athanasopoulos (2018). *Forecasting: principles and practice*. Melbourne, Australia: OTexts. <https://OTexts.org/fpp2>.
- Kahn, KB (1998). Revisiting top-down versus bottom-up forecasting. *The Journal of Business Forecasting* **17**(2), 14.
- Pennings, CL & J van Dalen (2017). Integrated hierarchical forecasting. *European Journal of Operational Research* **263**(2), 412–418.
- Tourism Research Australia (2005). Travel by Australians, September Quarter 2005. *Tourism Australia*.
- Wickramasuriya, SL, G Athanasopoulos & RJ Hyndman (2019). Optimal forecast reconciliation for hierarchical and grouped time series through trace minimization. *Journal of the American Statistical Association* **114**(526), 804–819.